

PatchReceipt

AI writes the patch.

PatchReceipt proves whether it deserves to ship.

A green test suite is not proof

THE GAP

AI patches can satisfy the reported example while failing retries, casing, mutation pressure, or scope policy. Reviewers need evidence that is reproducible, independent, and explainable.

Ordinary CI

- ✓ Existing tests pass
- ✗ Bug may not be reproduced
- ✗ No proof of test strength

PatchReceipt

- ✓ Proves before → after
- ✓ Sealed edge cases + PIT
- ✓ Scope and evidence receipt

One bug. Three plausible patches. Three defensible verdicts.

BUNDLED CHECKOUT BUG

Retried checkout requests must apply each case-insensitive coupon code once.

The verifier pack is sealed before candidates run.



REJECTED

distinct() passes the obvious case; 3 edge cases expose it.



PARTIAL


Correct result, but an unrelated production file changed.



VERIFIED

9/9 edge cases, 5/5 mutants, clean observed scope.

One ordered proof pipeline, not a trust score



Useful offline. Honest about its safety boundary.

NO API KEY

No runtime AI

JGit, Maven, JUnit, PIT and fixed verdict rules. No key. No credits.

Allowlisted hosting

Public API accepts only bundled caseld + patchld pairs.

Bounded execution

Single worker, queue cap, timeouts, memory and log limits.

Trusted local CLI

Refuses to execute unless the user passes an explicit safety flag.

Measured evidence from the six-patch corpus

MACHINE
READABLE

LOCAL EVALUATION

Two robust patches, one drifting patch, one edge-case trap, one build bypass, and one compile break.

6 / 6

expected verdicts

0

unsafe VERIFIED

25.9 s

final packaged run

Evidence before merge.

Deterministic verification

No runtime API required

Receipts judges can inspect